



Apache Spark Streaming Library

Spark Streaming

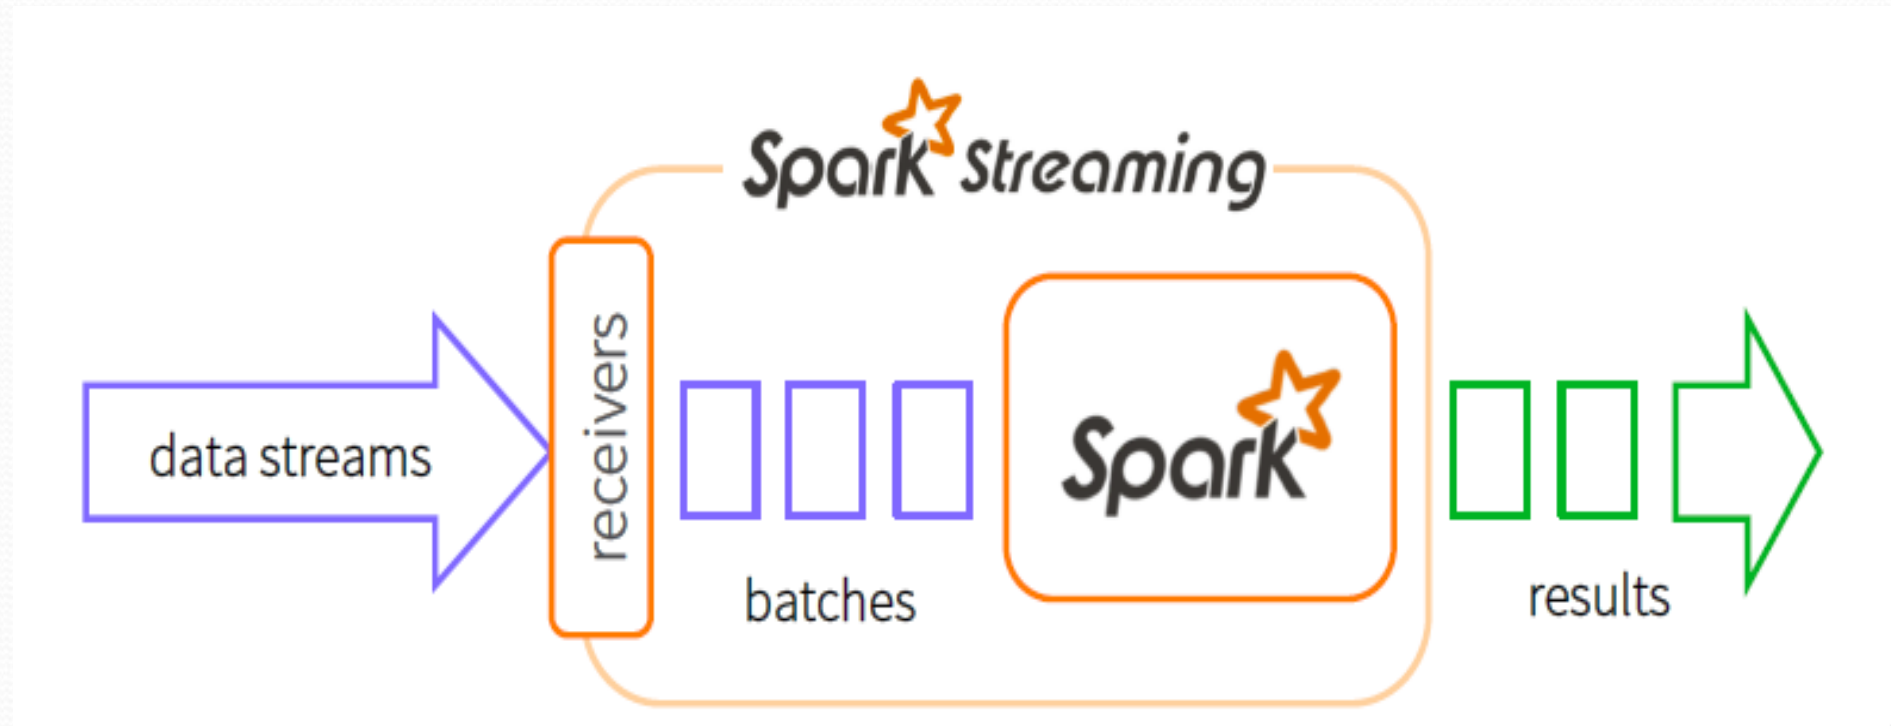


- Features of Spark Streaming
 - High Level API (joins, windows etc.)
 - Fault – Tolerant (exactly once semantics achievable)
 - Deep Integration with Spark Ecosystem (MLlib, SQL, GraphX etc.)



Spark Streaming

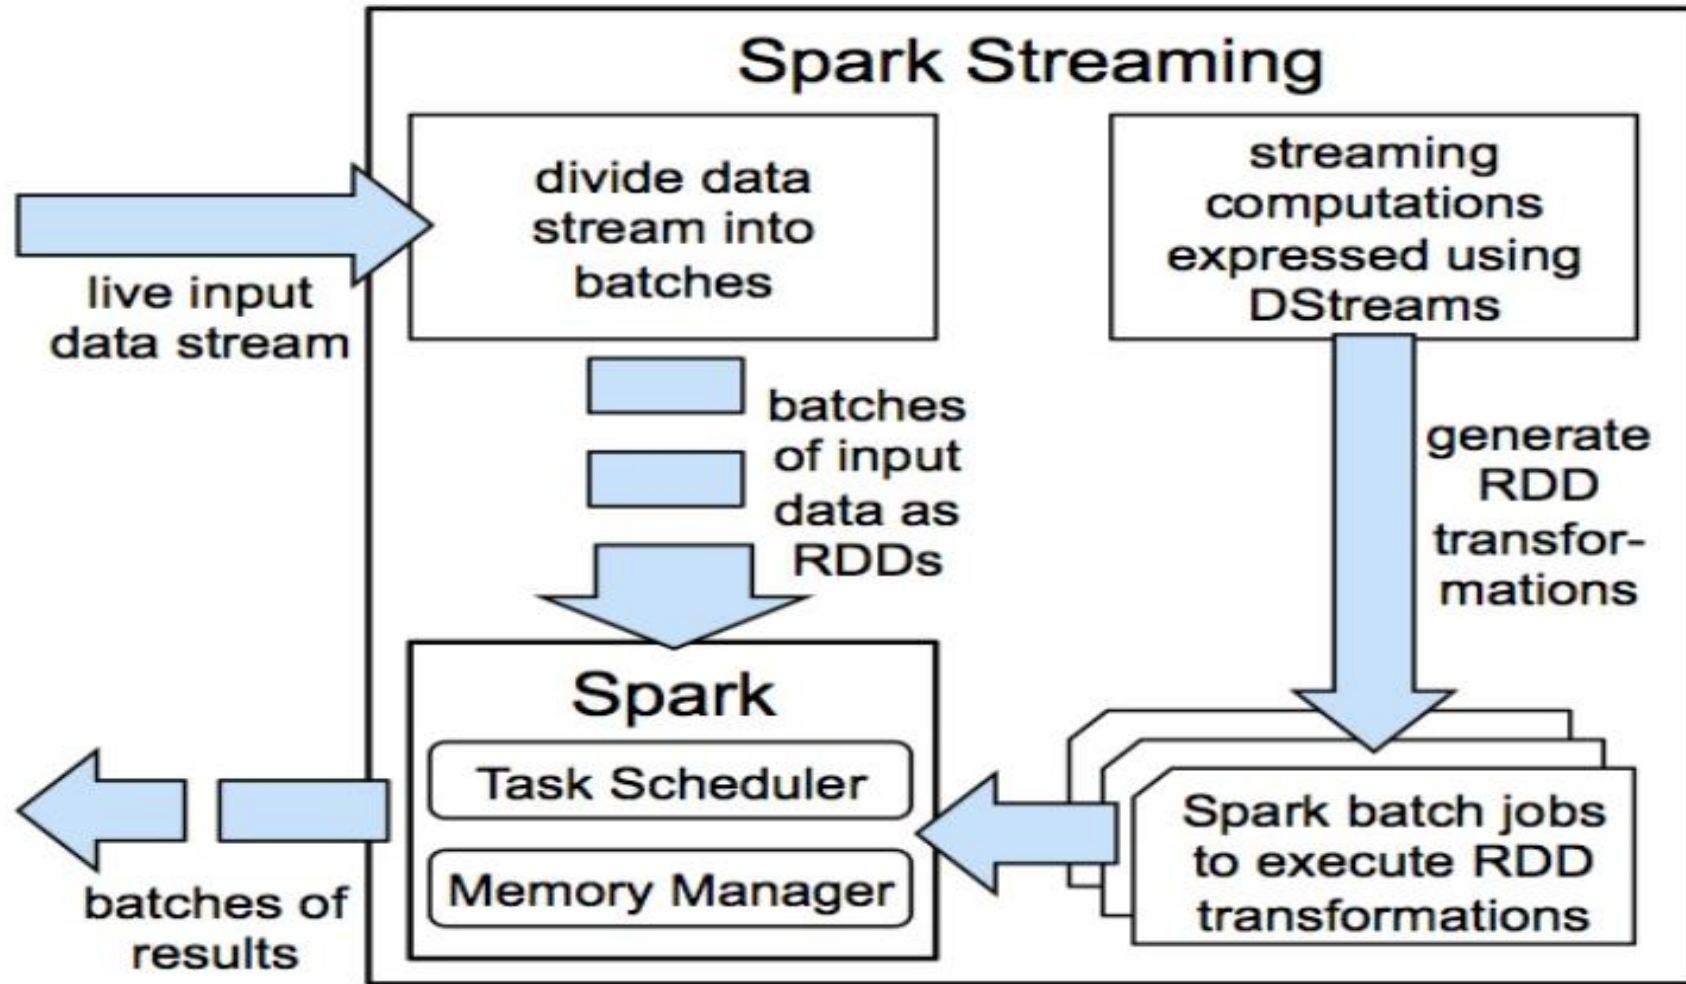
- Receivers receive data streams and chop them up into batches.
- Spark processes these batches and pushes out the results.






Architecture

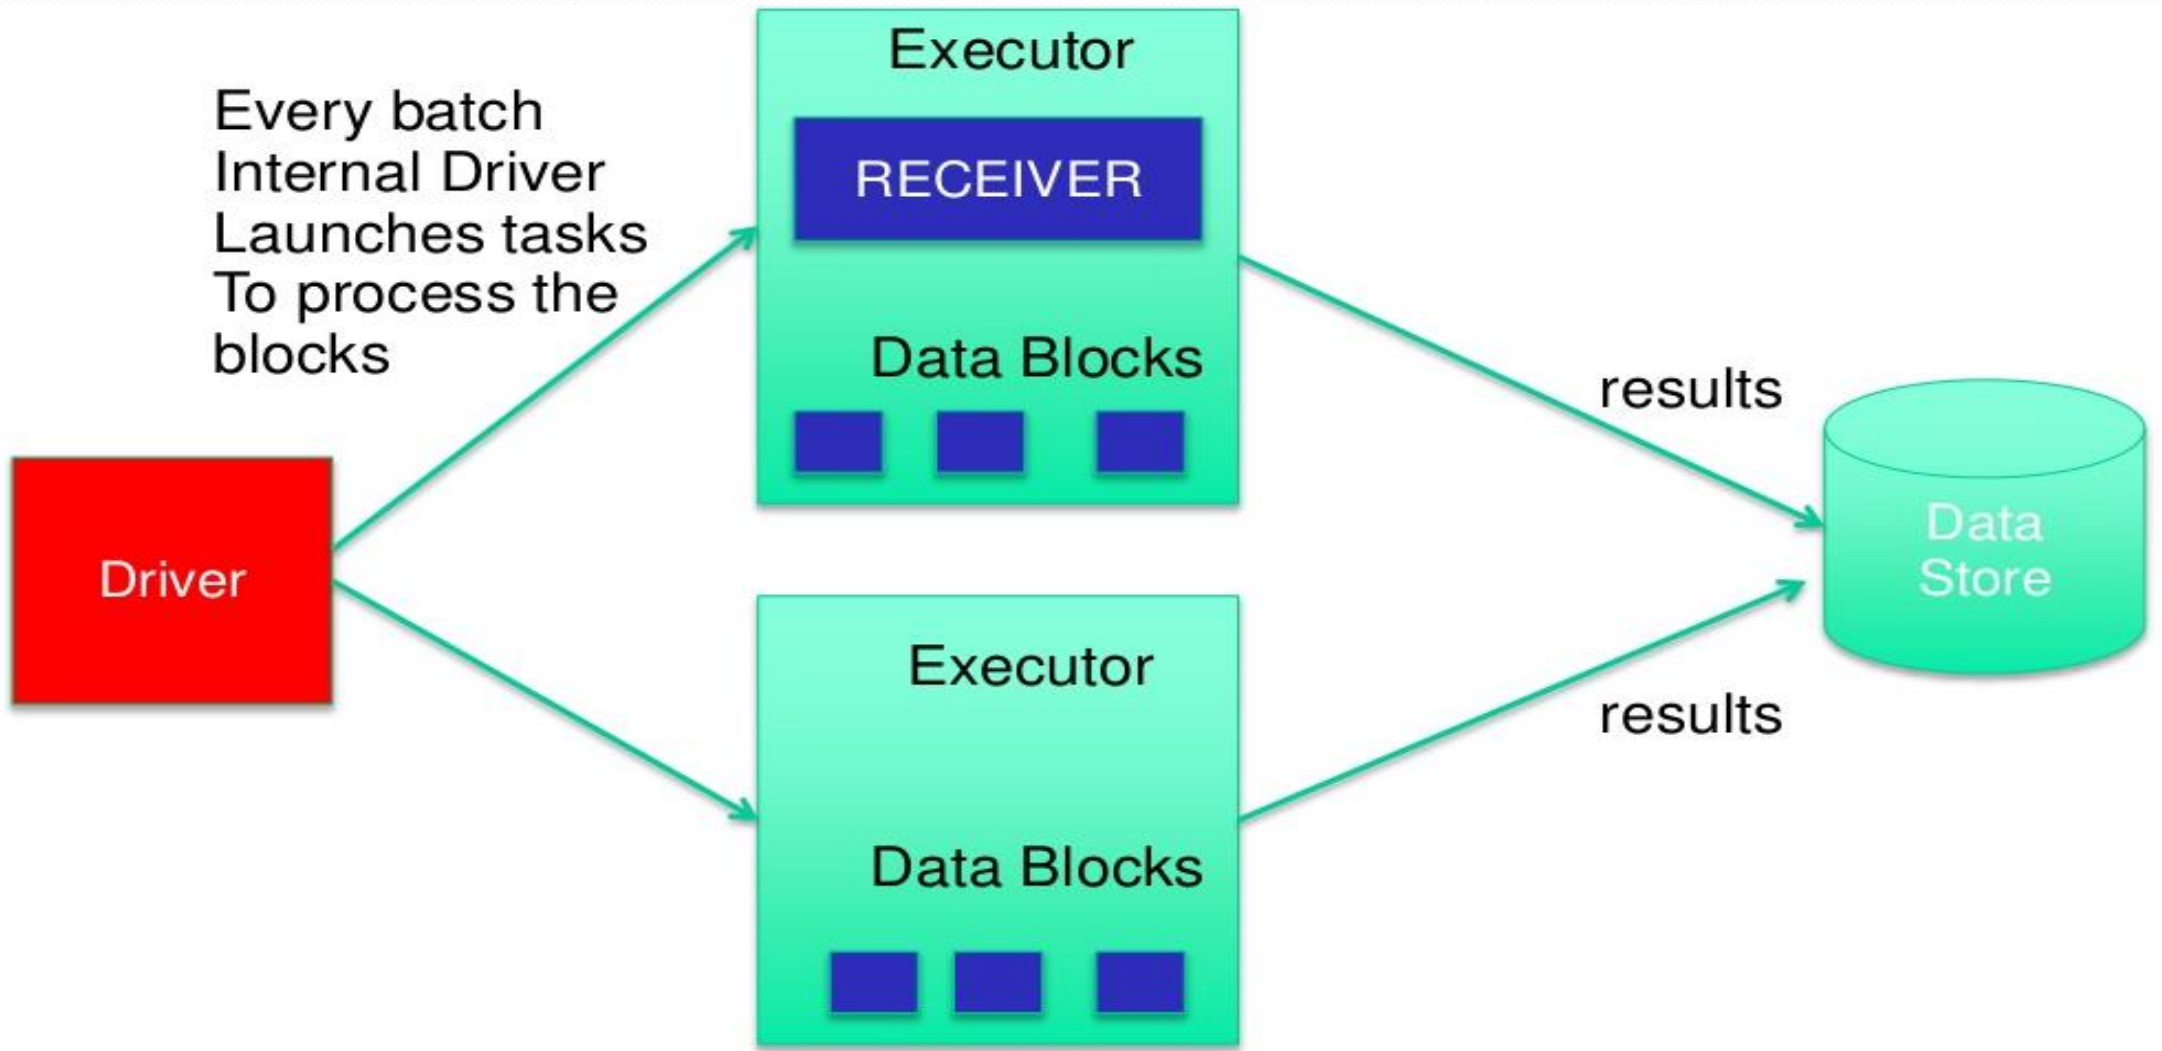
High Level Overview



Receiving Data



Processing Data



Word Count with Kafka

```
object WordCount {  
  def main(args: Array[String]) {  
    val context = new StreamingContext(new SparkConf(), Seconds(1))  
    val lines = KafkaUtils.createStream(context, ...)  
    val words = lines.flatMap(_.split(" "))  
    val wordCounts = words.map(x => (x,1)).reduceByKey(_ + _)  
    wordCounts.print()  
    context.start()  
    context.awaitTermination()  
  }  
}
```



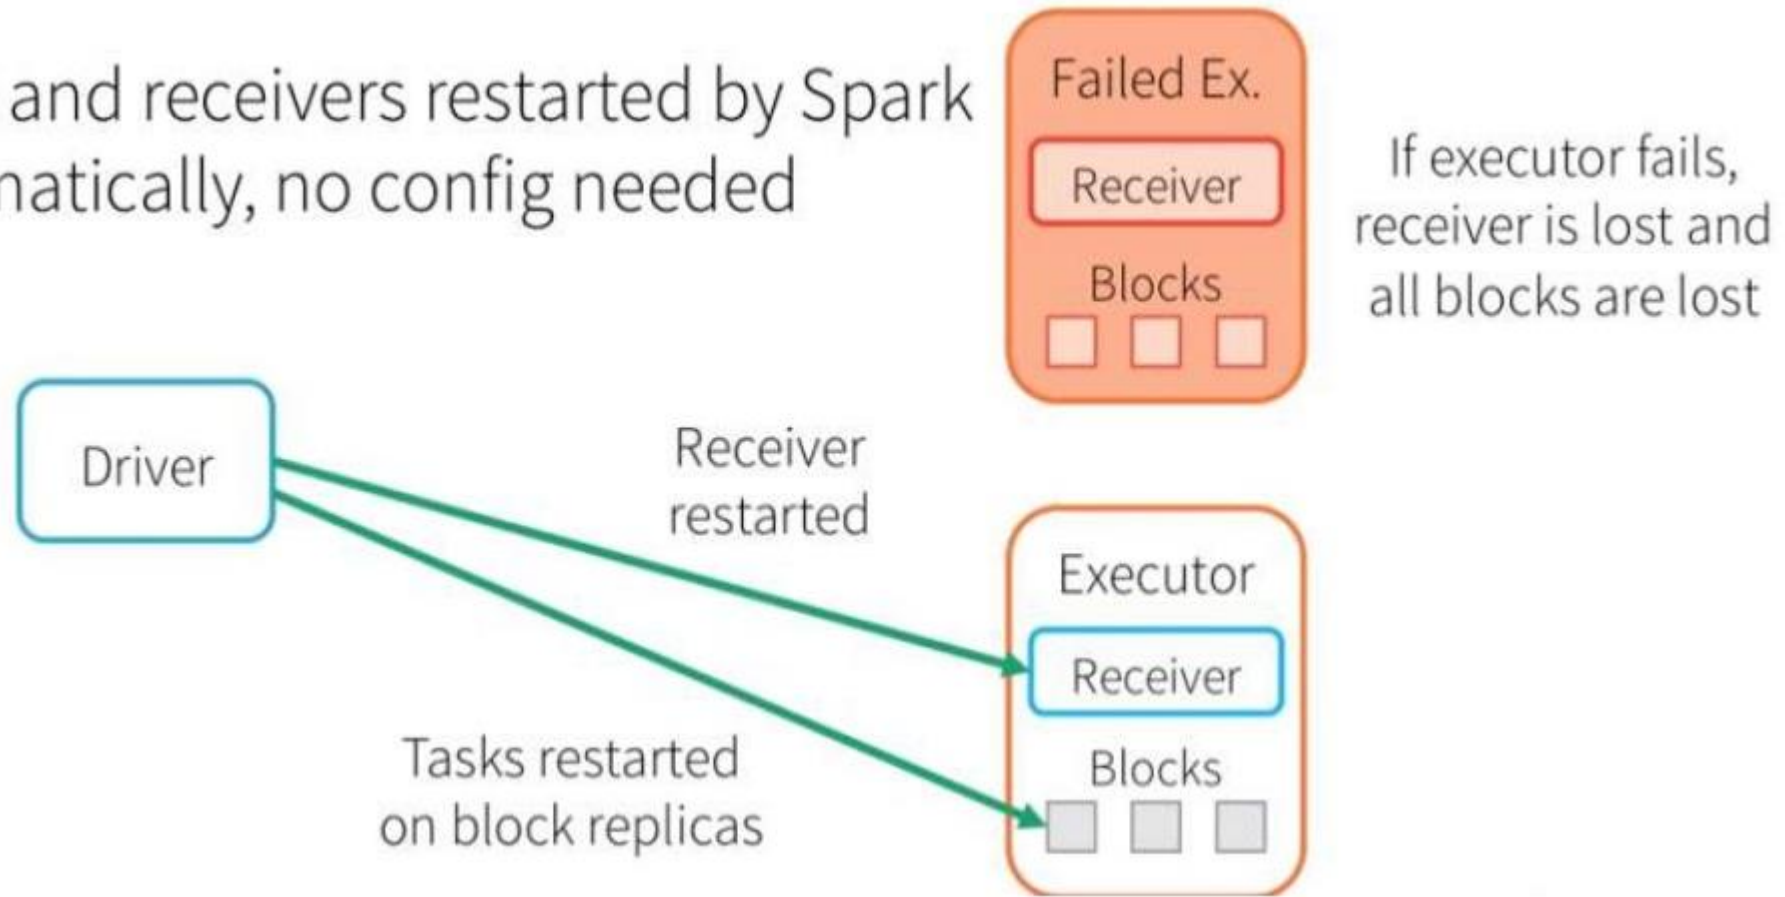

Fault Tolerance and Reliability

Fault Tolerance

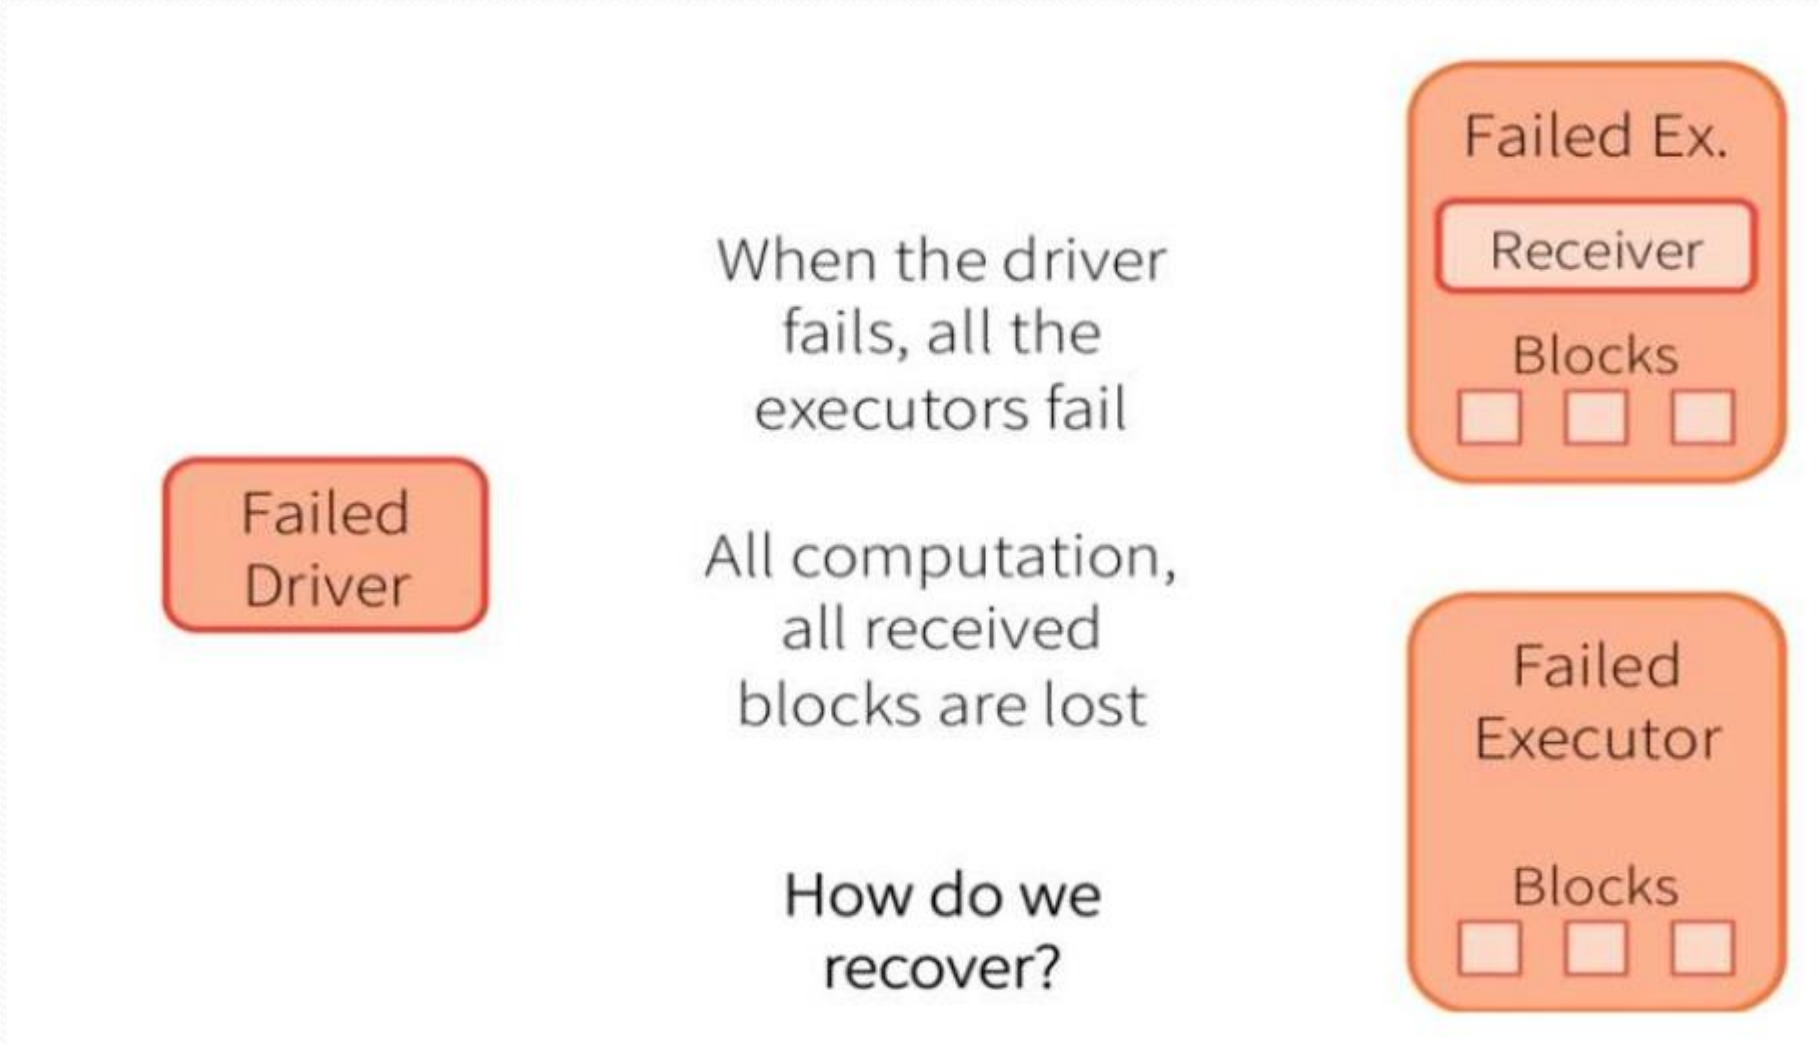
- Why Care ?
- Different guarantees for Data Loss
 - At least Once
 - Exactly Once
- What all can fail ?
 - Driver
 - Executor

What happens when Executor fails ?

Tasks and receivers restarted by Spark automatically, no config needed





What happens when Driver fails ?



Recovering Driver – DStream Checkpointing

DStream Checkpointing:

Periodically save the DAG of DStreams to fault-tolerant storage

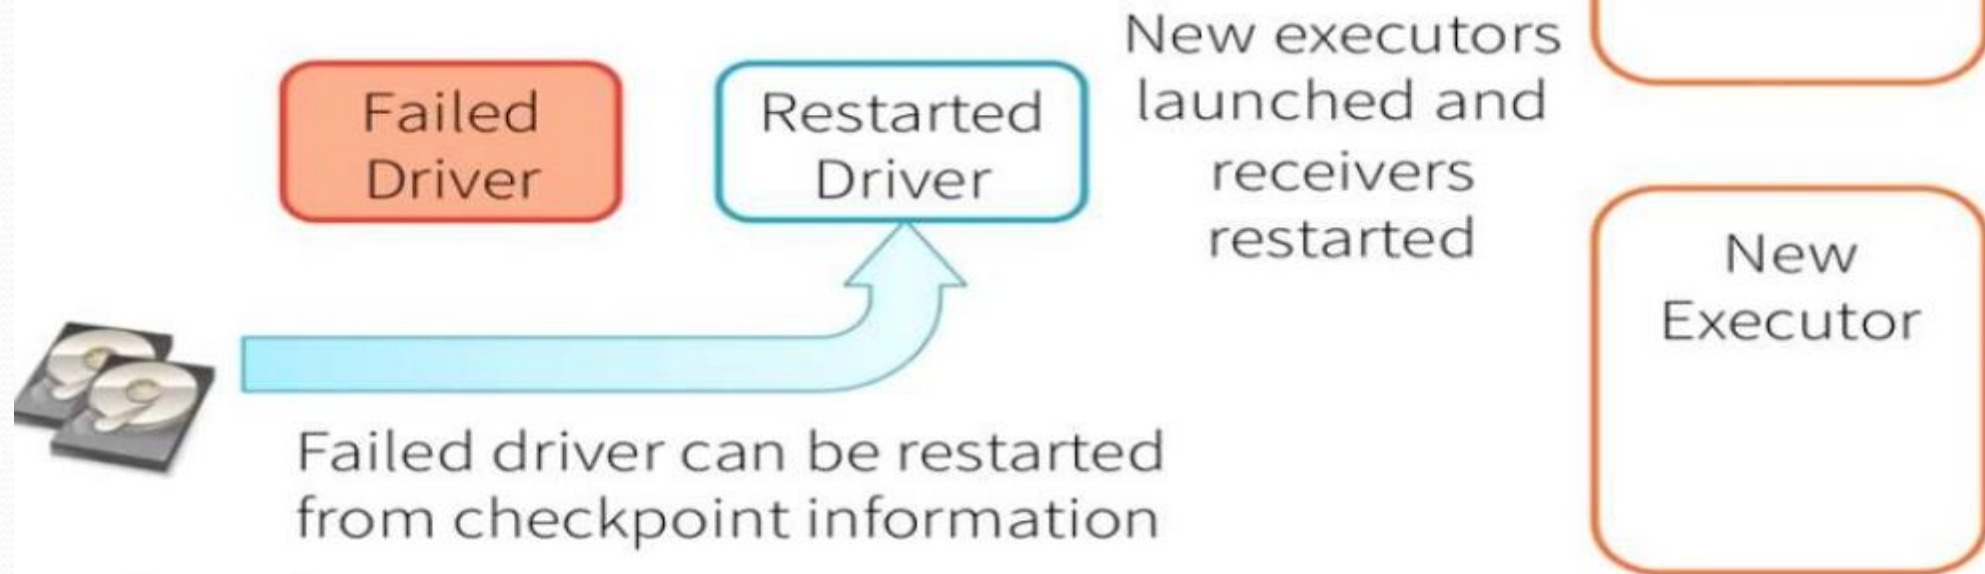


Recovering Driver – DStream Checkpointing

Driver restart

DStream Checkpointing:

Periodically save the DAG of DStreams to fault-tolerant storage



Driver restart – TODO List

- Configure automatic driver restart
 - Spark Standalone
 - YARN
- Set Checkpoint in HDFS compatible file system
streamingContext.checkpoint(hdfsDirectory)
- Ensure the Code uses checkpoints for recovery

```
Def setupStreamingContext() : StreamingContext = {  
    Val context = new StreamingContext(...)  
    Val lines = KafkaUtils.createStream(...)  
    ...  
    Context.checkpoint(hdfsDir)
```

```
Val context = StreamingContext.getOrCreate(hdfsDir, setupStreamingContext)  
Context.start()
```

Configuring Automatic Driver Restart

Spark Standalone – Use spark-submit with “cluster” mode and “--supervise”

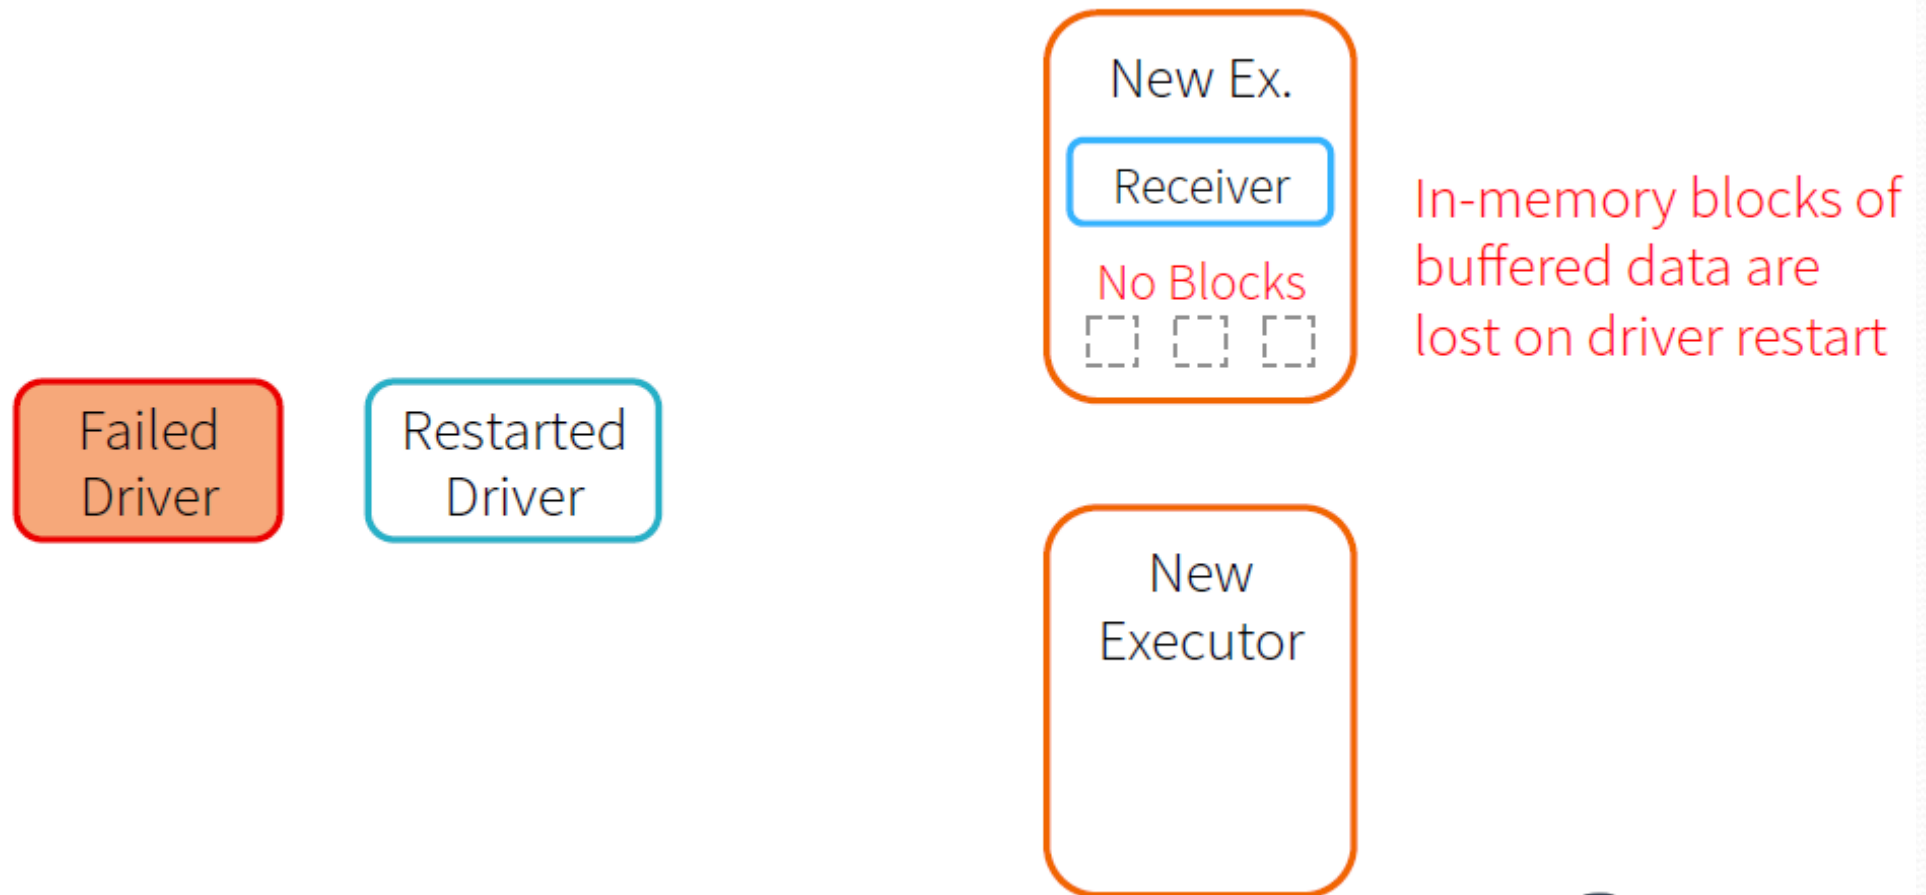
See <http://spark.apache.org/docs/latest/spark-standalone.html>

YARN – Use spark-submit in “cluster” mode

See YARN config “yarn.resourcemanager.am.max-attempts”


Mesos – Marathon can restart applications or use the “--supervise” flag.

Received blocks lost on Restart!



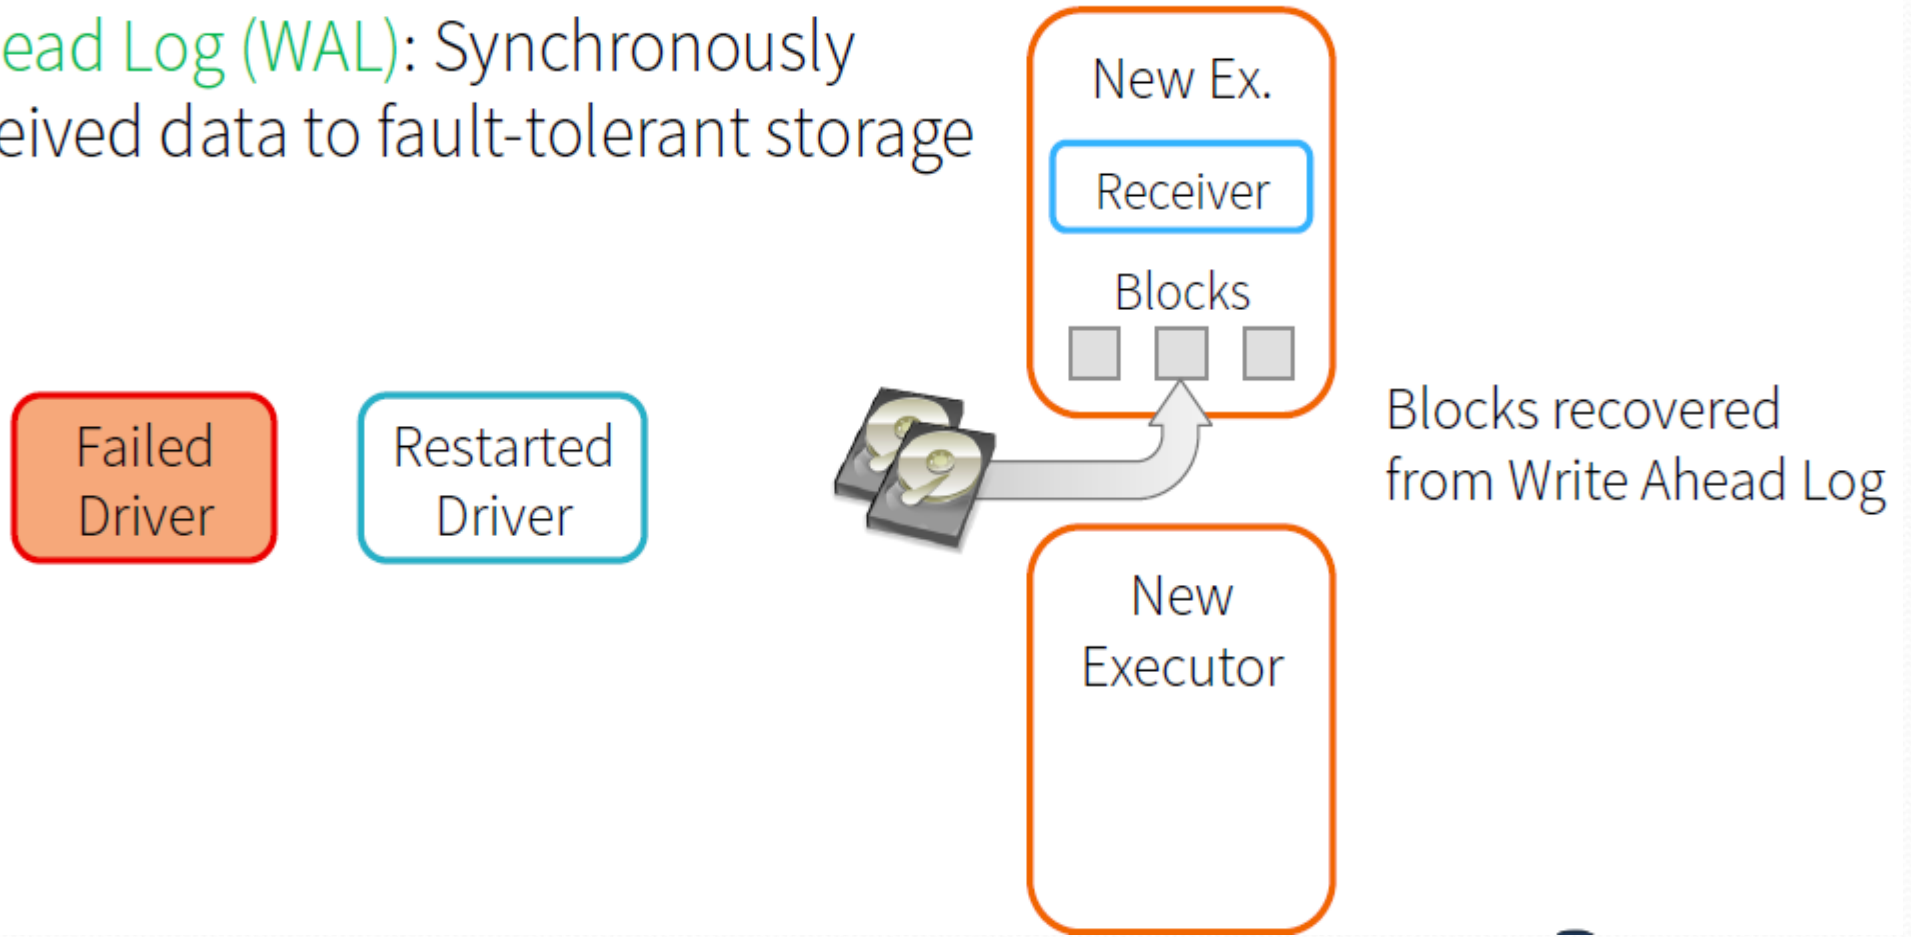
Recovering data with Write Ahead Logs

Write Ahead Log (WAL): Synchronously save received data to fault-tolerant storage



Recovering data with Write Ahead Logs

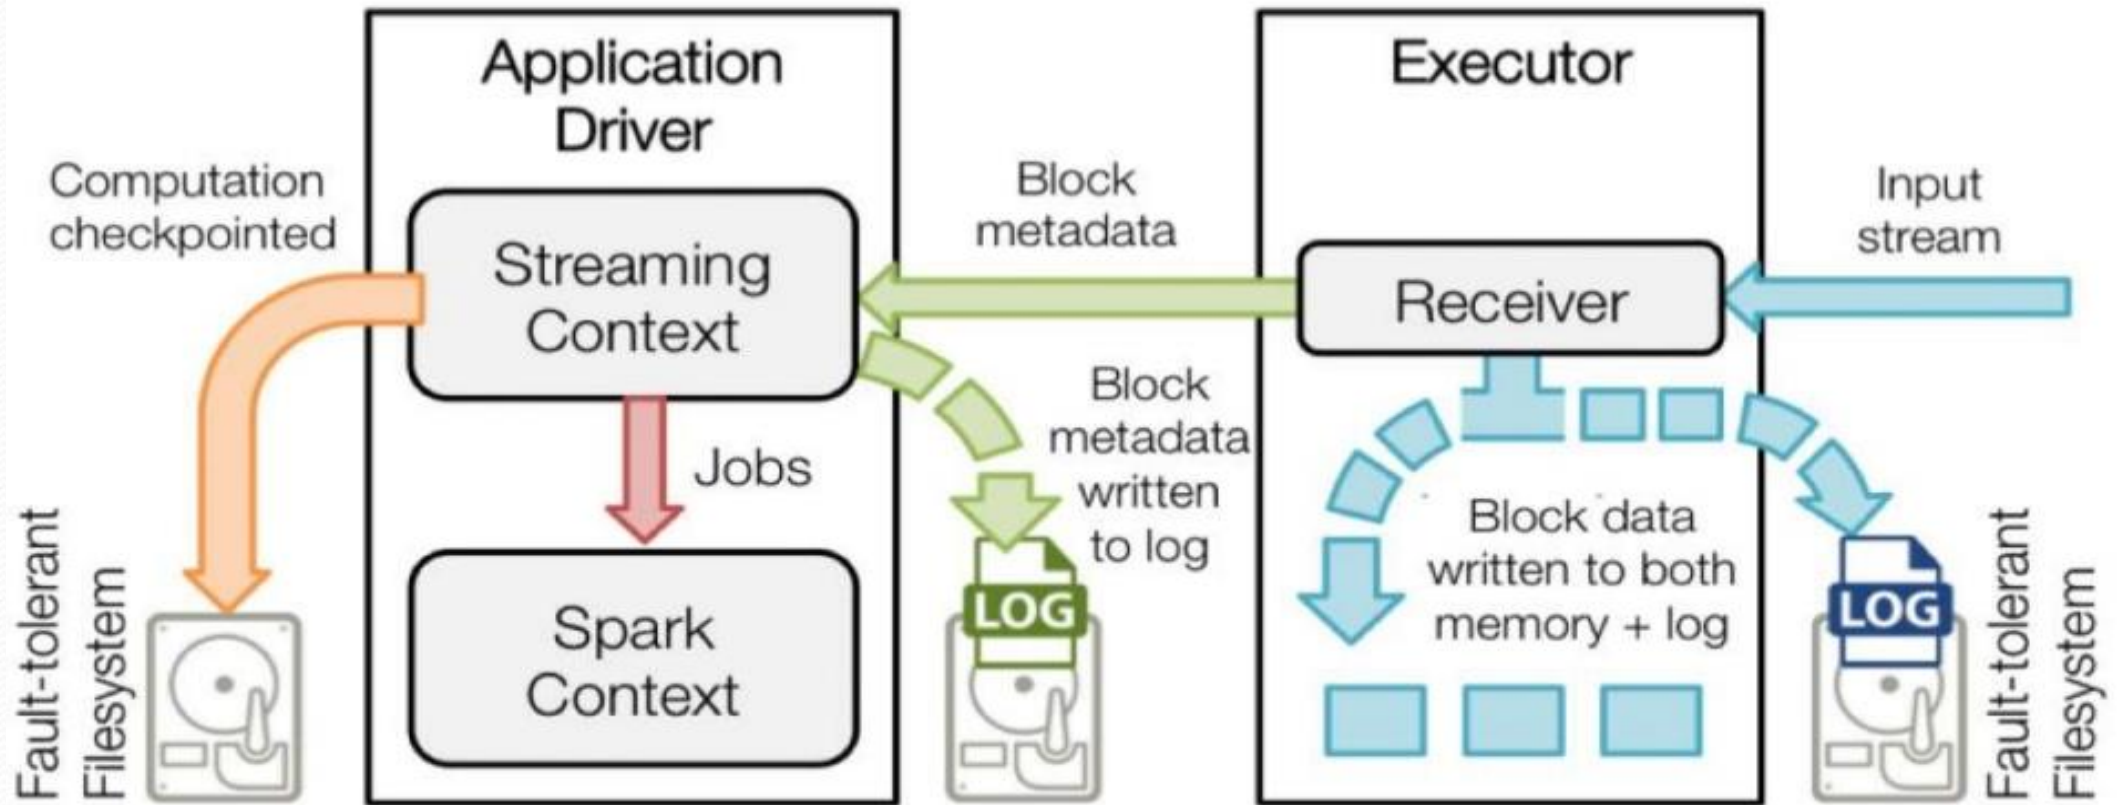
Write Ahead Log (WAL): Synchronously save received data to fault-tolerant storage



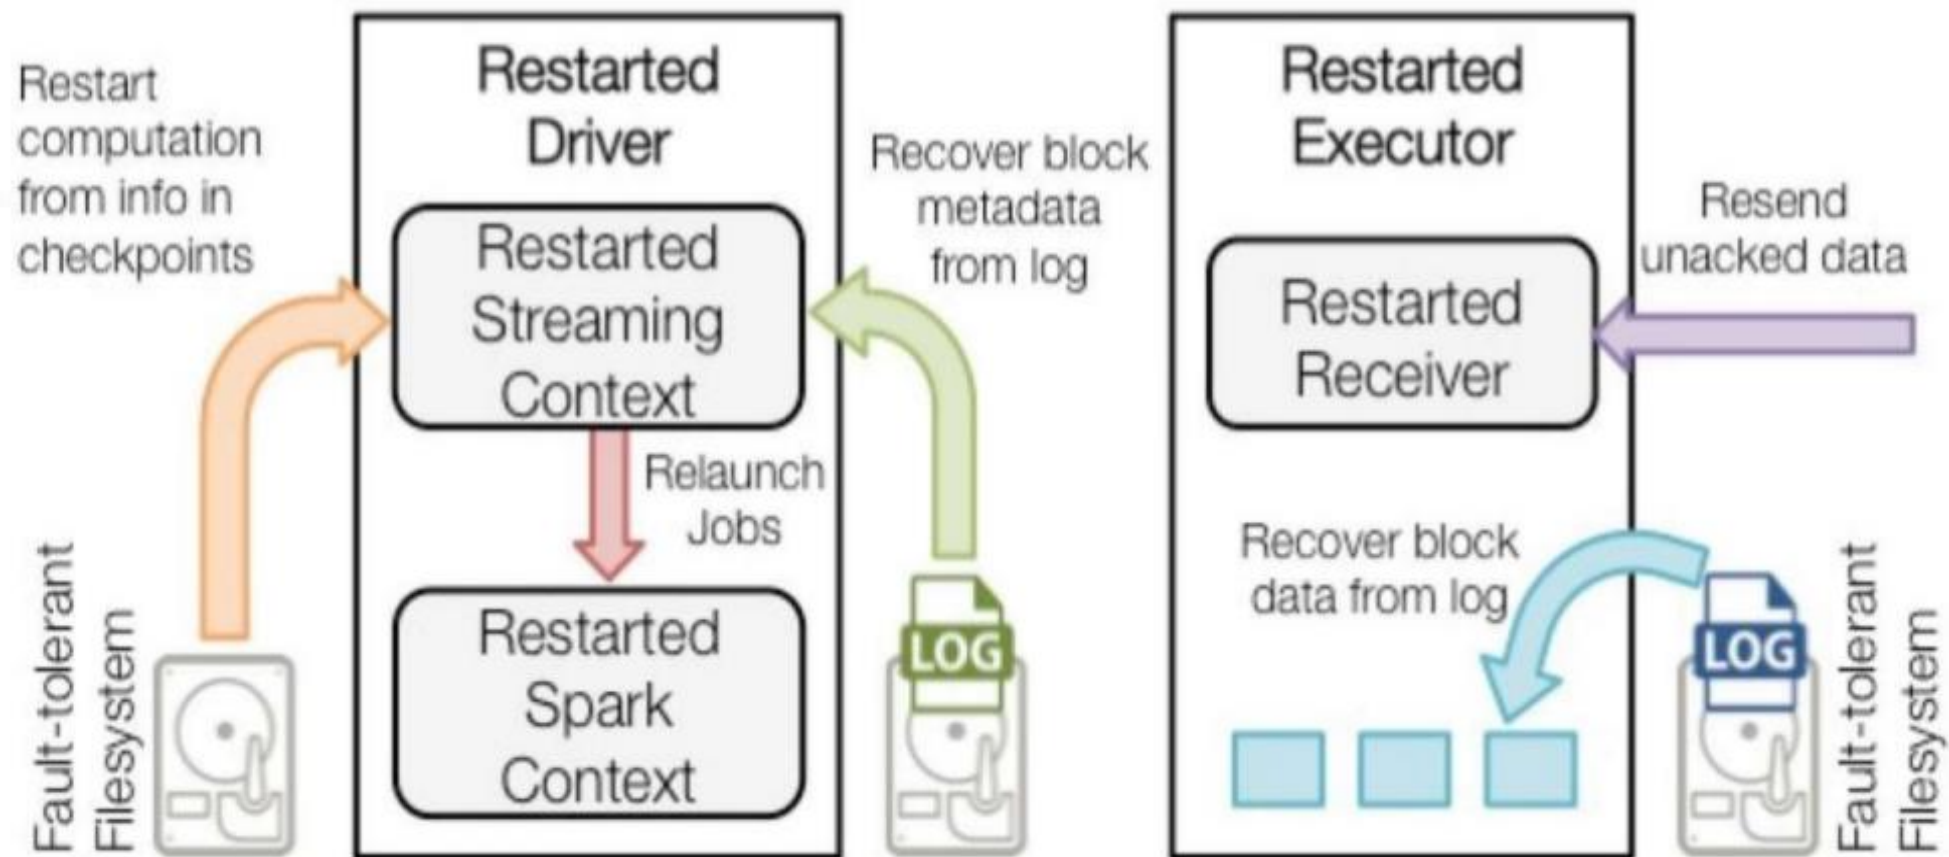
Configuration – Enabling WAL

1. Enable checkpointing, logs written in checkpoint directory
3. Enabled WAL in SparkConf configuration
`sparkConf.set("spark.streaming.receiver.writeAheadLog.enable", "true")`
3. Receiver should also be *reliable*
Acknowledge source only after data saved to WAL
Unacked data will be replayed from source by restarted receiver
5. Disable in-memory replication (already replicated by HDFS)

Normal Processing



Restarting Failed Driver



RDD Checkpointing

- Stateful stream processing can lead to long RDD lineages
- Long lineage = bad for fault-tolerance, too much recomputation
- RDD checkpointing saves RDD data to the fault-tolerant storage to limit lineage and recomputation